

## Detection Transformer (DETR) for Cervical Spine Fracture Detection (Kaggle Notebook)

- Add your private Kaggle dataset as an Input -->

```

1 # Import libraries
2 import os
3 import sys
4 import psutil
5 import h5py
6 import numpy as np
7 import matplotlib.pyplot as plt
8 import matplotlib.patches as patches
9 import warnings
10 import time
11 from torch.amp import autocast, GradScaler
12 import torch
13 import torchvision.transforms as T
14 from transformers import DetrForObjectDetection
15 from torch.optim import AdamW
16 from tqdm.auto import tqdm
17 import random
18 from sklearn.metrics import classification_report
19 import torchvision.ops as ops
20 from sklearn.metrics import f1_score
21 import cv2
22 from torch.utils.data import DataLoader
23
24 # Filter out PyTorch warning about meta parameters when loading model
25 warnings.filterwarnings("ignore", message=".*copying from a non-meta parameter.*")

```

```

1 # Get helper scripts from GitHub
2 !wget https://raw.githubusercontent.com/apmalinsky/AI-590-Capstone/refs/heads/main/scripts/dataloader.py
3 !wget https://raw.githubusercontent.com/apmalinsky/AI-590-Capstone/refs/heads/main/scripts/detection_batch_pr

--2025-12-01 23:38:36-- https://raw.githubusercontent.com/apmalinsky/AI-590-Capstone/refs/heads/main/scripts/dataloader.py
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133, 185.199.109.133, 185.199.111.133
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 4645 (4.5K) [text/plain]
Saving to: ‘dataloader.py’

dataloader.py      100%[=====] 4.54K --.-KB/s   in 0s

2025-12-01 23:38:36 (40.4 MB/s) - ‘dataloader.py’ saved [4645/4645]

--2025-12-01 23:38:36-- https://raw.githubusercontent.com/apmalinsky/AI-590-Capstone/refs/heads/main/scripts/detection\_batch\_prep.py
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.111.133, 185.199.108.133, 185.199.109.133
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.111.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 3039 (3.0K) [text/plain]
Saving to: ‘detection_batch_prep.py’

detection_batch_prep 100%[=====] 2.97K --.-KB/s   in 0s

2025-12-01 23:38:36 (45.0 MB/s) - ‘detection_batch_prep.py’ saved [3039/3039]

```

```

1 # Import helper scripts
2 from dataloader import get_dataloaders
3 from detection_batch_prep import prepare_batch

```

```

1 # View data file
2 data_path = '/kaggle/input/rsna-2022-hdf5-subset/fracture_dataset_subset.h5'
3
4 # Open the file
5 with h5py.File(data_path, 'r') as f:
6     # View key list, essentially our 'columns'
7     print(f'Keys: {list(f.keys())}')
8
9     # View column shapes
10    for key in f.keys():
11        dataset = f[key]
12        print(f'\nKey: {key}')
13        print(f'Shape: {dataset.shape}')
14        print(f'Dtype: {dataset.dtype}')

```

Keys: ['SliceNumber', 'StudyInstanceUID', 'bboxes', 'images', 'labels', 'split']

Key: SliceNumber  
Shape: (28812,)  
Dtype: object

Key: StudyInstanceUID  
Shape: (28812,)  
Dtype: object

Key: bboxes  
Shape: (28812, 10, 4)  
Dtype: float32

Key: images  
Shape: (28812, 256, 256)  
Dtype: float32

Key: labels  
Shape: (28812,)  
Dtype: int8

Key: split  
Shape: (28812,)  
Dtype: int8

```

1 # Get the number of physical CPU core
2 physical_cores = psutil.cpu_count(logical=False)
3 # Get the number of logical CPU cores (threads)
4 logical_cores = psutil.cpu_count(logical=True)
5
6 print(f'Physical CPU cores available: {physical_cores}')
7 print(f'Logical CPU cores (max workers) available: {logical_cores}')

```

Physical CPU cores available: 2  
Logical CPU cores (max workers) available: 4

```

1 # Configure parameters
2 batch_size = 100
3 num_workers = 2 # Increase workers to max available
4 learning_rate = 1e-4
5 backbone_lr = 1e-5
6 weight_decay = 1e-3 # We increase weight_decay to 1e-3 to prevent the training loss from dropping too fast (ov
7 num_epochs = 40
8 report_interval = 5      # Detailed report epoch frequency
9 patience = 20            # Stop early if no improvement for this many epochs
10 conf_threshold = 0.1    # Low threshold to monitor "potential" recall during training
11 image_size = 256
12 save_metric = 'f1'
13 device = 'cuda' if torch.cuda.is_available() else 'cpu' # for gpu acceleration
14 print(f'Using device: {device}')

```

Using device: cuda

```

1 # Apply data augmentation to prevent overfitting and memorization
2 # Applies safe augmentations, ones that do not move the bounding box

```

```

3 train_augs = T.Compose([
4     T.ToPILImage(),
5     T.ColorJitter(brightness=0.3, contrast=0.3), # Stronger lighting variation
6     T.GaussianBlur(kernel_size=3, sigma=(0.1, 2.0)), # Simulate blur
7     T.ToTensor(),
8 ])

```

```

1 # Get data loaders for detection
2 train_loader_det, val_loader_det, test_loader_det = get_dataloaders(
3     hdf5_path=data_path,
4     batch_size=batch_size,
5     task='detection',
6     num_workers=num_workers,
7     train_transform=train_augs
8 )

```

Loading dataset for task: 'detection'...  
 Loading splits...  
 DataLoaders created.

```

1 # Setup pre-trained DETR model
2 model = DetrForObjectDetection.from_pretrained(
3     "facebook/detr-resnet-50",
4     num_labels=1,
5     ignore_mismatched_sizes=True,
6     low_cpu_mem_usage=False
7 )
8 model.to(device)
9 print('DETR model loaded to device.')
10
11 # Freeze weights in early CNN layers to focus/speedup training for detection
12 for name, param in model.model.backbone.conv_encoder.model.named_parameters():
13     if "conv1" in name or "bn1" in name or "layer1" in name or "layer2" in name:
14         param.requires_grad = False
15 print('Frozen early CNN layers.')
16
17
18 # Set params and learning rates for optimizer
19 param_dicts = [
20     # Transformer params
21     {
22         "params": [p for n, p in model.named_parameters() if "backbone" not in n and p.requires_grad],
23         "lr": learning_rate,
24     },
25     # Backbone params
26     {
27         "params": [p for n, p in model.named_parameters() if "backbone" in n and p.requires_grad],
28         "lr": backbone_lr,
29     },
30 ]
31 optimizer = AdamW(param_dicts, lr=learning_rate, weight_decay=weight_decay)
32 print('Model initialized with balanced settings (anti-overfitting).')

```

Some weights of the model checkpoint at facebook/detr-resnet-50 were not used when initializing DetrForObjectDetection  
 - This IS expected if you are initializing DetrForObjectDetection from the checkpoint of a model trained on another  
 - This IS NOT expected if you are initializing DetrForObjectDetection from the checkpoint of a model that you expect  
 Some weights of DetrForObjectDetection were not initialized from the model checkpoint at facebook/detr-resnet-50 and  
 - class\_labels\_classifier.bias: found shape torch.Size([92]) in the checkpoint and torch.Size([2]) in the model ins  
 - class\_labels\_classifier.weight: found shape torch.Size([92, 256]) in the checkpoint and torch.Size([2, 256]) in t  
 You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.  
 DETR model loaded to device.  
 Froze early CNN layers.  
 Model initialized with balanced settings (anti-overfitting).

```

1 # Define helper function to process predictions by DETR model
2 def process_detr_predictions(outputs, conf_threshold=0.1, image_size=256):
3     # Get probability for "Fracture" class (index 0)
4     probas = outputs.logits.softmax(-1)[..., 0]
5     # Create mask for boxes above threshold

```

```

6     keep = probas > conf_threshold
7
8     # Convert relative [cx, cy, w, h] to absolute [x1, y1, x2, y2]
9     pred_boxes = outputs.pred_boxes.float()
10    cx, cy, w, h = pred_boxes.unbind(-1)
11    x1 = (cx - 0.5 * w) * image_size
12    y1 = (cy - 0.5 * h) * image_size
13    x2 = (cx + 0.5 * w) * image_size
14    y2 = (cy + 0.5 * h) * image_size
15    pred_xyxy = torch.stack([x1, y1, x2, y2], dim=-1)
16
17    return keep, pred_xyxy

```

```

1 # Define evaluation function for detection and classification metrics
2 def evaluate_dual_metrics(model, dataloader, device, conf_threshold=0.1, iou_threshold=0.5):
3     model.eval()
4     cls_y_true, cls_y_pred = [], [] # for classification
5     det_y_true, det_y_pred = [], [] # for detection
6
7     print(f'Starting Dual Evaluation (Conf: {conf_threshold}, IoU: {iou_threshold})...')
8
9     with torch.no_grad():
10         for images, targets in tqdm(dataloader, desc='Evaluating'):
11             pixel_values, _ = prepare_batch(images, targets, device, image_size=256, model_type='detr')
12             outputs = model(pixel_values=pixel_values)
13
14             # Process model predictions
15             keep, pred_xyxy = process_detr_predictions(outputs, conf_threshold)
16
17             # Loop through batch
18             for i in range(len(images)):
19                 gt_boxes = targets[i]['boxes'].to(device) # Get ground truth boxes
20                 img_pred_boxes = pred_xyxy[i][keep[i]] # Filter predictions for this image
21
22                 # Classification metrics
23                 has_fracture = len(gt_boxes) > 0 # True if patient has fracture, False if healthy
24                 pred_fracture = len(img_pred_boxes) > 0 # True if model predicts any fracture box
25                 cls_y_true.append(1 if has_fracture else 0)
26                 cls_y_pred.append(1 if pred_fracture else 0)
27
28                 # Detection metrics
29                 if len(gt_boxes) > 0: # Only evaluate detection if there is something to detect
30                     if len(img_pred_boxes) > 0:
31                         # For detection, calculate intersection over union (IoU) to compare bounding boxes
32                         # Calculate IoU between all ground truth and all predictions
33                         ious = ops.box_iou(gt_boxes, img_pred_boxes)
34
35                         # For each GT box, did we find a matching prediction?
36                         # We take the max IoU for each GT box
37                         max_ious_per_gt, _ = ious.max(dim=1)
38
39                         # A "Hit" is when the best overlap is > 0.5
40                         is_detected = max_ious_per_gt > iou_threshold
41
42                         det_y_true.extend([1] * len(gt_boxes)) # All GTs are "1" (should be found)
43                         det_y_pred.extend(is_detected.int().tolist())
44
45                 else:
46                     # We missed all boxes
47                     det_y_true.extend([1] * len(gt_boxes))
48                     det_y_pred.extend([0] * len(gt_boxes))
49
50                 # Print reports
51                 print('\n' + '='*60)
52                 print('CLASSIFICATION REPORT (Patient Diagnosis)')
53                 print(classification_report(cls_y_true, cls_y_pred, target_names=['Healthy', 'Fracture'], zero_division=0))
54
55                 print('\n' + '='*60)
56                 print('DETECTION REPORT (Localization)')
57                 print(classification_report(det_y_true, det_y_pred, target_names=['Missed Box', 'Found Box'], zero_divisio

```

```

1 # Check baseline (untrained) performance
2 print('Calculating Baseline Metrics ...')
3 evaluate_dual_metrics(model, val_loader_det, device)

Calculating Baseline Metrics ...
Starting Dual Evaluation (Conf: 0.1, IoU: 0.5)...
Evaluating:  0%|          | 0/42 [00:00<?, ?it/s]

=====
CLASSIFICATION REPORT (Patient Diagnosis)
precision    recall   f1-score   support
Healthy       0.00     0.00     0.00      3135
Fracture      0.25     1.00     0.40     1045

accuracy           0.25      4180
macro avg        0.12     0.50     0.20      4180
weighted avg     0.06     0.25     0.10      4180

=====
DETECTION REPORT (Localization)
precision    recall   f1-score   support
Missed Box    0.00     0.00     0.00       0
Found Box     1.00     0.11     0.20     1045

accuracy           0.11      1045
macro avg        0.50     0.06     0.10      1045
weighted avg     1.00     0.11     0.20      1045

```

```

1 # Define validation function to be used during training
2 def run_validation(model, dataloader, device, conf_threshold=0.1, iou_threshold=0.5):
3     model.eval()
4     total_loss = 0.0
5     true_positives = 0
6     false_positives = 0
7     false_negatives = 0
8
9     with torch.no_grad():
10         for images, targets in tqdm(dataloader, desc='Validating', leave=False):
11             pixel_values, labels = prepare_batch(images, targets, device, image_size=256, model_type='detr')
12             with torch.amp.autocast('cuda'):
13                 outputs = model(pixel_values=pixel_values, labels=labels)
14                 if outputs.loss is not None:
15                     total_loss += outputs.loss.item()
16
17             # Process model predictions
18             keep, pred_xyxy = process_detr_predictions(outputs, conf_threshold)
19
20             for i in range(len(images)):
21                 gt_boxes = targets[i]['boxes'].to(device)
22                 pred_boxes = pred_xyxy[i][keep[i]]                      # Filter by confidence
23
24                 if len(gt_boxes) > 0:
25                     # Patient has fracture
26                     if len(pred_boxes) > 0:
27                         # Check overlap (IoU)
28                         ious = ops.box_iou(gt_boxes, pred_boxes)
29
30                         # If any predicted box overlaps any gt box by > 0.5, it's a hit
31                         if (ious.max() > iou_threshold).item():
32                             true_positives += 1
33                         else:
34                             # Predicted a box, but in the wrong place
35                             false_negatives += 1
36                             false_positives += 1 # Penalize for bad box
37                         else:
38                             # Missed completely

```

```

39         false_negatives += 1
40     else:
41         # Healthy Patient
42         if len(pred_boxes) > 0:
43             false_positives += 1 # False alarm
44
45     # Calculate scores
46     avg_loss = total_loss / len(dataloader)
47
48     precision = true_positives / (true_positives + false_positives + 1e-6)
49     recall = true_positives / (true_positives + false_negatives + 1e-6)
50     f1_score = 2 * (precision * recall) / (precision + recall + 1e-6)
51
52     return avg_loss, f1_score, recall, precision

```

```

1 # Training loop, optimizing for best f1-score
2 scaler = GradScaler('cuda')
3
4 # Tracking metric history
5 best_f1 = 0.0
6 early_stopping_counter = 0
7 history = {'train_loss': [], 'val_loss': [], 'f1': [], 'recall': [], 'precision': []}
8
9 print(f'Starting Training Loop (Target: Best F1-Score, {conf_threshold} conf)...')
10
11 start_time = time.time()
12
13 for epoch in range(num_epochs):
14     # Training Phase
15     model.train()
16     train_loss_accum = 0.0
17     pbar = tqdm(train_loader_det, desc=f'Epoch {epoch+1}/{num_epochs} [Train]')
18
19     for images, targets in pbar:
20         pixel_values, labels = prepare_batch(images, targets, device, image_size=256, model_type='detr')
21         optimizer.zero_grad(set_to_none=True)
22
23         with autocast(device_type='cuda', dtype=torch.float16):
24             outputs = model(pixel_values=pixel_values, labels=labels)
25             loss = outputs.loss
26
27             scaler.scale(loss).backward()
28             scaler.unscale_(optimizer)
29             torch.nn.utils.clip_grad_norm_(model.parameters(), 0.1)
30             scaler.step(optimizer)
31             scaler.update()
32
33             train_loss_accum += loss.item()
34             pbar.set_postfix({'loss': f'{loss.item():.4f}'})
35
36             avg_train_loss = train_loss_accum / len(train_loader_det)
37             history['train_loss'].append(avg_train_loss)
38
39     # Validation Phase
40     avg_val_loss, val_f1, val_recall, val_prec = run_validation(
41         model, val_loader_det, device, conf_threshold=conf_threshold
42     )
43
44     history['val_loss'].append(avg_val_loss)
45     history['f1'].append(val_f1)
46     history['recall'].append(val_recall)
47     history['precision'].append(val_prec)
48
49     print(f' -> Loss: Train={avg_train_loss:.4f} Val={avg_val_loss:.4f}')
50     print(f' -> Metrics: F1={val_f1:.2%} | Recall={val_recall:.2%} | Prec={val_prec:.2%}')
51
52     # Saving logic (best F1)
53     if val_f1 > best_f1:
54         print(f' [*] NEW BEST F1! ({best_f1:.2%} -> {val_f1:.2%}). Saving "detr_best_f1.pth"...')

```

```
55     best_f1 = val_f1
56     best_recall_at_f1 = val_recall
57     torch.save(model.state_dict(), 'detr_best_f1.pth')
58     early_stopping_counter = 0
59 else:
60     early_stopping_counter += 1
61     print(f'  [!] No improvement. Patience: {early_stopping_counter}/{patience}')
62
63 if early_stopping_counter >= patience:
64     print('\n[STOP] Early stopping triggered.')
65     break
66
67 total_time = (time.time() - start_time) / 60
68 print(f'\nTraining Finished in {total_time:.1f} minutes.'
```



```
Starting Training Loop (Target: Best F1-Score, 0.1 conf)...
Epoch 1/40 [Train]:  0% | 0/208 [00:00<?, ?it/s]
Exception ignored in: <function tqdm._del_ at 0x7cc8219fae80>
Traceback (most recent call last):
  File "/usr/local/lib/python3.11/dist-packages/tqdm/std.py", line 1148, in __del__
    self.close()
  File "/usr/local/lib/python3.11/dist-packages/tqdm/notebook.py", line 279, in close
    self.disp(bar_style='danger', check_delay=False)
    ^^^^^^^^^^
AttributeError: 'tqdm' object has no attribute 'disp'
Validating:  0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.8797 Val=1.6076
-> Metrics: F1=0.00% | Recall=0.00% | Prec=0.00%
[!] No improvement. Patience: 1/20
Epoch 2/40 [Train]:  0% | 0/208 [00:00<?, ?it/s]
Exception ignored in: <function tqdm._del_ at 0x7cc8219fae80>
Traceback (most recent call last):
  File "/usr/local/lib/python3.11/dist-packages/tqdm/std.py", line 1148, in __del__
    self.close()
  File "/usr/local/lib/python3.11/dist-packages/tqdm/notebook.py", line 279, in close
    self.disp(bar_style='danger', check_delay=False)
    ^^^^^^^^^^
AttributeError: 'tqdm' object has no attribute 'disp'
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x7cc85fda8360>
Traceback (most recent call last):
  File "/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py", line 1618, in __del__
    self._shutdown_workers()
  File "/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py", line 1601, in _shutdown_workers
    if w.is_alive():
    ^^^^^^^^^^^^
  File "/usr/lib/python3.11/multiprocessing/process.py", line 160, in is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child process'
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x7cc85fda8360>
Traceback (most recent call last):
  File "/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py", line 1618, in __del__
    self._shutdown_workers()
  File "/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py", line 1601, in _shutdown_workers
    if w.is_alive():
    ^^^^^^^^^^^^
  File "/usr/lib/python3.11/multiprocessing/process.py", line 160, in is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child process'
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function tqdm._del_ at 0x7cc8219fae80>
Traceback (most recent call last):
  File "/usr/local/lib/python3.11/dist-packages/tqdm/std.py", line 1148, in __del__
    self.close()
  File "/usr/local/lib/python3.11/dist-packages/tqdm/notebook.py", line 279, in close
    self.disp(bar_style='danger', check_delay=False)
    ^^^^^^^^^^
AttributeError: 'tqdm' object has no attribute 'disp'
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x7cc85fda8360>
Traceback (most recent call last):
  File "/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py", line 1618, in __del__
    self._shutdown_workers()
  File "/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py", line 1601, in _shutdown_workers
    if w.is_alive():
    ^^^^^^^^^^^^
  File "/usr/lib/python3.11/multiprocessing/process.py", line 160, in is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child process'
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at 0x7cc85fda8360>
Traceback (most recent call last):
  File "/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py", line 1618, in __del__
    self._shutdown_workers()
  File "/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py", line 1601, in _shutdown_workers
    if w.is_alive():
    ^^^^^^^^^^^^
  File "/usr/lib/python3.11/multiprocessing/process.py", line 160, in is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child process'
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
AssertionError: can only test a child process
Validating:  0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.6470 Val=1.5740
```

```

-> Metrics: F1=0.00% | Recall=0.00% | Prec=0.00%
[!] No improvement. Patience: 2/20
Epoch 3/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.4988 Val=1.5397
-> Metrics: F1=1.67% | Recall=0.86% | Prec=25.71%
[*] NEW BEST F1! (0.00% -> 1.67%). Saving "detr_best_f1.pth"...
Epoch 4/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.4267 Val=1.5234
-> Metrics: F1=0.00% | Recall=0.00% | Prec=0.00%
[!] No improvement. Patience: 1/20
Epoch 5/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.3415 Val=1.5506
-> Metrics: F1=0.00% | Recall=0.00% | Prec=0.00%
[!] No improvement. Patience: 2/20
Epoch 6/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.2832 Val=1.4956
-> Metrics: F1=1.89% | Recall=1.05% | Prec=9.17%
[*] NEW BEST F1! (1.67% -> 1.89%). Saving "detr_best_f1.pth"...
Epoch 7/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.2236 Val=1.3836
-> Metrics: F1=11.41% | Recall=11.67% | Prec=11.15%
[*] NEW BEST F1! (1.89% -> 11.41%). Saving "detr_best_f1.pth"...
Epoch 8/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.1900 Val=1.3781
-> Metrics: F1=13.59% | Recall=13.68% | Prec=13.49%
[*] NEW BEST F1! (11.41% -> 13.59%). Saving "detr_best_f1.pth"...
Epoch 9/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.1442 Val=1.4520
-> Metrics: F1=5.20% | Recall=2.97% | Prec=21.09%
[!] No improvement. Patience: 1/20
Epoch 10/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.1164 Val=1.4615
-> Metrics: F1=10.21% | Recall=7.56% | Prec=15.71%
[!] No improvement. Patience: 2/20
Epoch 11/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.0812 Val=1.4014
-> Metrics: F1=15.50% | Recall=16.08% | Prec=14.96%
[*] NEW BEST F1! (10.21% -> 15.50%). Saving "detr_best_f1.pth"...
Epoch 12/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.0575 Val=1.3823
-> Metrics: F1=13.81% | Recall=11.20% | Prec=18.03%
[!] No improvement. Patience: 1/20
Epoch 13/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.0359 Val=1.4117
-> Metrics: F1=5.88% | Recall=3.54% | Prec=17.29%
[!] No improvement. Patience: 2/20
Epoch 14/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.0426 Val=1.4315
-> Metrics: F1=13.42% | Recall=9.28% | Prec=24.19%
[!] No improvement. Patience: 3/20
Epoch 15/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=1.0019 Val=1.4287
-> Metrics: F1=11.10% | Recall=11.20% | Prec=11.01%
[!] No improvement. Patience: 4/20
Epoch 16/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.9780 Val=1.4881
-> Metrics: F1=18.23% | Recall=18.09% | Prec=18.37%
[*] NEW BEST F1! (11.10% -> 18.23%). Saving "detr_best_f1.pth"...
Epoch 17/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.9592 Val=1.4585
-> Metrics: F1=5.92% | Recall=3.16% | Prec=47.14%
[!] No improvement. Patience: 1/20

```

```

Epoch 18/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.9366 Val=1.4801
-> Metrics: F1=15.90% | Recall=11.58% | Prec=25.37%
[!] No improvement. Patience: 2/20
Epoch 19/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.9245 Val=1.4489
-> Metrics: F1=14.96% | Recall=11.39% | Prec=21.79%
[!] No improvement. Patience: 3/20
Epoch 20/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.9103 Val=1.4995
-> Metrics: F1=16.12% | Recall=12.63% | Prec=22.26%
[!] No improvement. Patience: 4/20
Epoch 21/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.8827 Val=1.4380
-> Metrics: F1=14.09% | Recall=15.41% | Prec=12.98%
[!] No improvement. Patience: 5/20
Epoch 22/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.8743 Val=1.4841
-> Metrics: F1=14.90% | Recall=19.81% | Prec=11.94%
[!] No improvement. Patience: 6/20
Epoch 23/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.8714 Val=1.4971
-> Metrics: F1=17.06% | Recall=21.44% | Prec=14.17%
[!] No improvement. Patience: 7/20
Epoch 24/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.8455 Val=1.4877
-> Metrics: F1=15.33% | Recall=13.78% | Prec=17.27%
[!] No improvement. Patience: 8/20
Epoch 25/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.8459 Val=1.5768
-> Metrics: F1=8.41% | Recall=5.74% | Prec=15.71%
[!] No improvement. Patience: 9/20
Epoch 26/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.8215 Val=1.5003
-> Metrics: F1=14.95% | Recall=17.80% | Prec=12.88%
[!] No improvement. Patience: 10/20
Epoch 27/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.8138 Val=1.4740
-> Metrics: F1=13.90% | Recall=10.14% | Prec=22.08%
[!] No improvement. Patience: 11/20
Epoch 28/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.8042 Val=1.4864
-> Metrics: F1=13.15% | Recall=13.68% | Prec=12.65%
[!] No improvement. Patience: 12/20
Epoch 29/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.8006 Val=1.4206
-> Metrics: F1=13.32% | Recall=14.83% | Prec=12.09%
[!] No improvement. Patience: 13/20
Epoch 30/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.7984 Val=1.4111
-> Metrics: F1=14.56% | Recall=16.94% | Prec=12.76%
[!] No improvement. Patience: 14/20
Epoch 31/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.7765 Val=1.5949
-> Metrics: F1=13.17% | Recall=11.39% | Prec=15.62%
[!] No improvement. Patience: 15/20
Epoch 32/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.7859 Val=1.5160
-> Metrics: F1=11.79% | Recall=12.54% | Prec=11.12%
[!] No improvement. Patience: 16/20
Epoch 33/40 [Train]: 0% | 0/208 [00:00<?, ?it/s]
Validating: 0% | 0/42 [00:00<?, ?it/s]
-> Loss: Train=0.7707 Val=1.5200

```

```

1 # Get training and validation metric history for plots
2 train_losses = history['train_loss']
3 val_losses = history['val_loss']
4 val_f1s = history['f1']
5 val_recalls = history['recall']
6 val_precisions = history['precision']
7
8 # Find the best epoch (Max f1 score)
9 best_idx = np.argmax(val_f1s)
10 best_epoch = best_idx + 1
11
12 t_loss = train_losses[:best_idx+1]
13 v_loss = val_losses[:best_idx+1]
14 v_f1 = val_f1s[:best_idx+1]
15 v_recall = val_recalls[:best_idx+1]
16 v_precision = val_precisions[:best_idx+1]
17 epochs = range(1, best_epoch + 1)
18
19 print(f'Best Validation f1: {val_f1s[best_idx]:.4f} at Epoch {best_epoch}')

```


[...] NO improvement. patience: 5/20  
 BestValF1:0.10% at Epoch 340/208 [00:00<?, ?it/s]  
 Validating: 0% | 0/42 [00.00<?, ?it/s]

```

1 # Plot training and validation loss until best epoch
2 plt.figure(figsize=(10, 6))
3 plt.plot(epochs, t_loss, label='Training Loss', marker='o', color='blue')
4 plt.plot(epochs, v_loss, label='Validation Loss', marker='x', color='red')
5 plt.axvline(x=best_epoch, color='green', linestyle='--', label='Best Model Saved')
6 plt.title(f'Training and Validation Loss (Truncated at Best Epoch {best_epoch})')
7 plt.xlabel('Epoch')
8 plt.ylabel('Loss')
9 plt.legend()
10 plt.grid(True)
11 plt.show()

```

Training and Validation Loss (Truncated at Best Epoch 34)




```

1 # Plot validation report metrics (f1-score, recall, precision)
2 plt.figure(figsize=(10, 6))
3 plt.plot(epochs, v_f1, label='f1 Score', marker='+', color='green')
4 plt.plot(epochs, v_recall, label='Validation Recall', marker='o', color='blue')
5 plt.plot(epochs, v_precision, label='Validation Precision', marker='x', color='red')

```

```
6 plt.axvline(x=best_epoch, color='green', linestyle='--', label='Best Model Saved')
7 plt.title(f'Validation Report Metrics During Training (Truncated at Best Epoch {best_epoch})')
8 plt.xlabel('Epoch')
9 plt.ylabel('Score')
10 plt.legend()
11 plt.grid(True)
12 plt.show()
```



```
1 # Load the best saved checkpoint
2 checkpoint = torch.load('detr_best_f1.pth', map_location=device)
3 model.load_state_dict(checkpoint)
4 model.to(device)
5 print('Loaded best model from saved checkpoint.')
6
7 # Evaluate on Test Set
8 evaluate_dual_metrics(model, test_loader_det, device, conf_threshold=0.1)
```

```

Loaded best model from saved checkpoint.
Starting Dual Evaluation (Conf: 0.1, IoU: 0.5)...
Evaluating:  0%|          | 0/39 [00:00<?, ?it/s]

=====
CLASSIFICATION REPORT (Patient Diagnosis)
precision    recall   f1-score   support

Healthy       0.83     0.66     0.74     2916
Fracture      0.37     0.60     0.46     972

accuracy           0.65     3888
macro avg        0.60     0.63     0.60     3888
weighted avg     0.72     0.65     0.67     3888

=====
DETECTION REPORT (Localization)
precision    recall   f1-score   support

Missed Box    0.00     0.00     0.00      0
Found Box     1.00     0.27     0.42     972

accuracy           0.27     972
macro avg        0.50     0.13     0.21     972
weighted avg     1.00     0.27     0.42     972

```

## Optimization

```

1 def find_best_threshold(model, dataloader, device, start=0.01, end=0.5, step=0.02):
2     model.eval()
3
4     # Store raw data to avoid re-running model
5     all_pred_probs = []
6     all_pred_boxes = []
7     all_gt_boxes = []
8
9     with torch.no_grad():
10         print('Getting Predictions on Validation Set (Once)...')
11         for images, targets in tqdm(dataloader, desc='Collecting Predictions'):
12             pixel_values, _ = prepare_batch(images, targets, device, image_size=256, model_type='detr')
13             outputs = model(pixel_values=pixel_values)
14
15             # Get Probabilities
16             probas = outputs.logits.softmax(-1)[..., 0]
17
18             # Get and Scale Boxes
19             pred_boxes = outputs.pred_boxes
20             cx, cy, w, h = pred_boxes.unbind(-1)
21             x1 = (cx - 0.5 * w) * 256
22             y1 = (cy - 0.5 * h) * 256
23             x2 = (cx + 0.5 * w) * 256
24             y2 = (cy + 0.5 * h) * 256
25             scaled_boxes = torch.stack([x1, y1, x2, y2], dim=-1)
26
27             # Store ground truths and predictions
28             for i in range(len(images)):
29                 all_pred_probs.append(probas[i].cpu())
30                 all_pred_boxes.append(scaled_boxes[i].cpu())
31
32             # Handle GT boxes
33             gt_box = targets[i]['boxes'].cpu()
34             all_gt_boxes.append(gt_box)
35
36     print('\nTesting Thresholds...')
37     best_f1 = 0.0
38     best_thresh = 0.0

```

```

39
40     # Loop through thresholds
41     thresholds = np.arange(start, end, step)
42
43     for thresh in thresholds:
44         y_true_det = []
45         y_pred_det = []
46
47         for i in range(len(all_pred_probs)):
48             gt = all_gt_boxes[i]
49             probs = all_pred_probs[i]
50             boxes = all_pred_boxes[i]
51
52             # Apply threshold
53             keep_idxs = probs > thresh
54             pred_b = boxes[keep_idxs]
55
56             # Check if Patient has Fracture (GT > 0)
57             if len(gt) > 0 and gt.shape[-1] == 4:
58                 # Check if Model Predicted Fracture
59                 if len(pred_b) > 0 and pred_b.shape[-1] == 4:
60                     ious = ops.box_iou(gt, pred_b)
61                     max_ious, _ = ious.max(dim=1)
62
63                     # Hit if overlap > 0.5
64                     is_hit = (max_ious > 0.5).any().item()
65
66                     y_true_det.append(1)
67                     y_pred_det.append(1 if is_hit else 0)
68                 else:
69                     # Missed (No prediction above threshold)
70                     y_true_det.append(1)
71                     y_pred_det.append(0)
72             else:
73                 # Healthy Patient (GT=0)
74                 if len(pred_b) > 0:
75                     # False Alarm
76                     y_true_det.append(0)
77                     y_pred_det.append(1)
78                 else:
79                     # True Negative (Correctly ignored)
80                     y_true_det.append(0)
81                     y_pred_det.append(0)
82
83             # Calculate F1 for this threshold
84             curr_f1 = f1_score(y_true_det, y_pred_det, zero_division=0)
85
86             if curr_f1 > best_f1:
87                 best_f1 = curr_f1
88                 best_thresh = thresh
89
90     print(f'\n[*] BEST RESULT:')
91     print(f'Best Threshold: {best_thresh:.2f}')
92     print(f'Best F1-Score: {best_f1:.4f}')
93
94     return best_thresh

```

```

1 # Tune confidence threshold on validation set
2 print('Tuning Threshold on Validation Set...')
3 optimal_threshold = find_best_threshold(model, val_loader_det, device)
4
5 print(f'\nThe optimal confidence threshold is: {optimal_threshold:.2f}')

```

```
Tuning Threshold on Validation Set...
Getting Predictions on Validation Set (Once)...
Collecting Predictions:  0%|          | 0/42 [00:00<?, ?it/s]

Testing Thresholds...

[*] BEST RESULT:
Best Threshold: 0.01
Best F1-Score:  0.3174

The optimal confidence threshold is: 0.01
```

```
1 # Final evaluation on test set with optimal confidence threshold
2 print('Final Evaluation on Test Set...')
3 evaluate_dual_metrics(model, test_loader_det, device, conf_threshold=optimal_threshold)
```

```
Final Evaluation on Test Set...
Starting Dual Evaluation (Conf: 0.01, IoU: 0.5)...
Evaluating:  0%|          | 0/39 [00:00<?, ?it/s]
```

```
=====
CLASSIFICATION REPORT (Patient Diagnosis)
precision    recall   f1-score   support
Healthy       0.98     0.25      0.40      2916
Fracture      0.31     0.99      0.47      972
accuracy      0.44
macro avg     0.64     0.62      0.43      3888
weighted avg  0.81     0.44      0.42      3888
```

```
=====
DETECTION REPORT (Localization)
precision    recall   f1-score   support
Missed Box    0.00     0.00      0.00       0
Found Box     1.00     0.61      0.75      972
accuracy      0.61
macro avg     0.50     0.30      0.38      972
weighted avg  1.00     0.61      0.75      972
```

```
1 # Visualize fracture truth/predictions
2 def visualize_fracture_predictions(model, dataloader, device, conf_threshold=0.01, num_images=4):
3     model.eval()
4     print(f'Scanning for {num_images} Fracture examples (Threshold: {conf_threshold})...')
5
6     # Collect only fracture examples
7     fracture_imgs = []
8     fracture_targets = []
9
10    # Iterate through dataloader until we have enough
11    for batch_images, batch_targets in dataloader:
12        for img, tgt in zip(batch_images, batch_targets):
13            if len(tgt['boxes']) > 0: # Check if there is a Ground Truth box
14                fracture_imgs.append(img)
15                fracture_targets.append(tgt)
16
17            if len(fracture_imgs) >= num_images:
18                break
19            if len(fracture_imgs) >= num_images:
20                break
21
22    print(f'Found {len(fracture_imgs)} fracture images. Running Inference... ')
23
24    # Run model for bounding box predictions
25    pixel_values, _ = prepare_batch(fracture_imgs, fracture_targets, device, image_size=256, model_type='det'
26
```

```

27     with torch.no_grad():
28         outputs = model(pixel_values=pixel_values)
29
30         # Get probabilities
31         probas = outputs.logits.softmax(-1)[..., 0]
32         # Keep boxes above threshold
33         keep = probas > conf_threshold
34
35         # Process boxes
36         pred_boxes = outputs.pred_boxes
37         cx, cy, w, h = pred_boxes.unbind(-1)
38         x1 = (cx - 0.5 * w) * 256
39         y1 = (cy - 0.5 * h) * 256
40         x2 = (cx + 0.5 * w) * 256
41         y2 = (cy + 0.5 * h) * 256
42         w_px = w * 256
43         h_px = h * 256
44         pred_xywh = torch.stack([x1, y1, w_px, h_px], dim=-1)
45
46         # Plot images with bounding boxes
47         num_cols = 2
48         num_rows = len(fracture_imgs)
49         fig, axes = plt.subplots(num_rows, num_cols, figsize=(10, 4 * num_rows))
50         if num_rows == 1: axes = [axes]
51
52         for i in range(num_rows):
53             # Left for ground truth
54             ax_gt = axes[i][0] if num_rows > 1 else axes[0]
55
56             img_tensor = fracture_imgs[i]
57             if img_tensor.shape[0] == 3:
58                 img_np = img_tensor.permute(1, 2, 0)
59             else:
60                 img_np = img_tensor.squeeze(0)
61
62             ax_gt.imshow(img_np, cmap='gray')
63             ax_gt.set_title(f'Ground Truth (Fracture)')
64             ax_gt.axis('off')
65
66             # Draw GT Boxes (Green)
67             for box in fracture_targets[i]['boxes']:
68                 box_data = box.tolist()
69                 if len(box_data) >= 4:
70                     x1, y1, x2, y2 = box_data[:4]
71                     w = x2 - x1
72                     h = y2 - y1
73                     rect = patches.Rectangle((x1, y1), w, h, linewidth=2, edgecolor='lime', facecolor='none')
74                     ax_gt.add_patch(rect)
75
76             # Right for model prediction
77             ax_pred = axes[i][1] if num_rows > 1 else axes[1]
78             ax_pred.imshow(img_np, cmap='gray')
79             ax_pred.set_title(f'Model Prediction (Conf > {conf_threshold})')
80             ax_pred.axis('off')
81
82             # Find the single best score for this image
83             best_score, best_idx = probas[i].max(dim=0)
84
85             # Only display if it passes the minimum safety threshold
86             if best_score > conf_threshold:
87                 # Get the specific box
88                 top_box = pred_xywh[i][best_idx]
89
90                 # Extract coordinates
91                 box_data = top_box.tolist()
92                 x, y, w, h = box_data
93
94                 # Draw
95                 rect = patches.Rectangle((x, y), w, h, linewidth=2, edgecolor='red', facecolor='none')
96                 ax_pred.add_patch(rect)

```